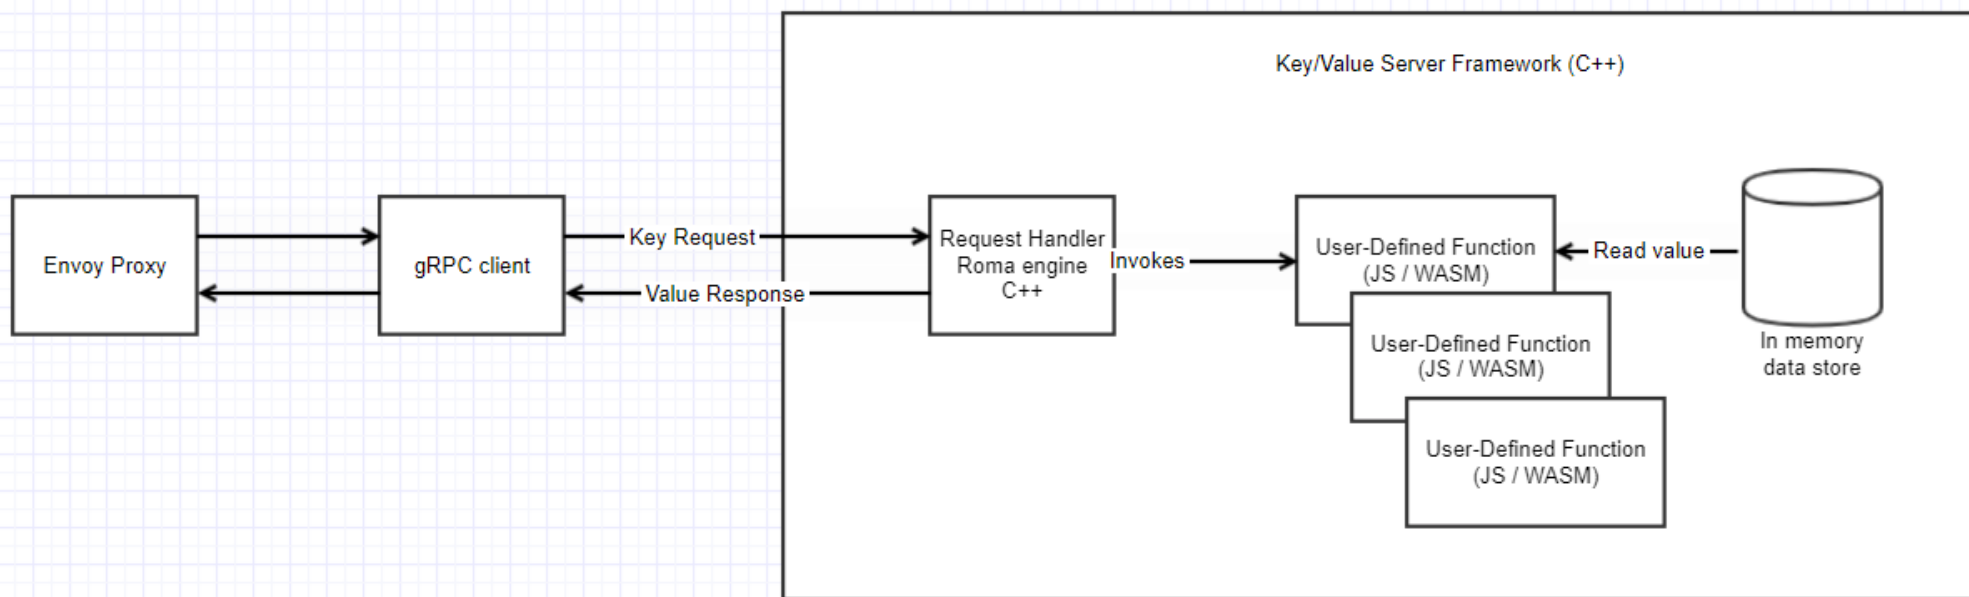


Key value server / ROMA benchmarks

Fabian Höring

f.horing@criteo.com

Non TEE server setup on our own infra



Server setup

Deploy locally compiled key-value server on 1 instance with 16 cores & 16 GB of memory with a very basic implementation

- No batching
- it reads the keys from the request
- it queries the in-memory datastore to get the values of those keys
- it replies with those values
- C# asp.net vs Google key/value service with same input & output

Web load test with gatling

- Traffic dump with calls to the key/value server (keys 1 to 1000)
- QPS 500-100000
- Assert results for server payload

Results ASP.NET vs KV server

Service	mean response time	90 percentile	99 percentile	max QPS	max Memory
c# fledge bidding	3ms	5ms	9ms	45k	2.5 GB
c# fledge bidding (at 10k QPS)	1ms	2ms	2ms	10k	2.5 GB
key value service	6ms	8ms	9ms	5k	1 GB

Conclusion ASP.NET vs KV server

- Google implementation can handle 1000s of queries per second with ms second latency
- Asp.net can handle 10 times more queries (45k QPS vs 5k QPS) with better mean response times (3ms vs 6ms)

WASM with c#/c++

- Emscripten is an LLVM/Clang-based compiler that compiles C and C++ source code to WebAssembly

```
async function getModule() {  
    var Module = {  
        instantiateWasm: function (imports, successCallback) {  
            var module = new WebAssembly.Module(wasm_array);  
            var instance = new WebAssembly.Instance(module, imports);  
            Module.testWasmInstantiationSucceeded = 1;  
            successCallback(instance);  
            return instance.exports;  
        },  
    };  
    return await wasmModule(Module);  
}
```

WASM with dotnet (dotnet 9 preview)

dotnet workload install wasi-experimental

```
<Project Sdk="Microsoft.NET.Sdk">
  <PropertyGroup>
    <TargetFramework>net9.0</TargetFramework>
    <RuntimeIdentifier>wasi-wasm</RuntimeIdentifier>
    <OutputType>Exe</OutputType>
    <PublishTrimmed>true</PublishTrimmed>
    <WasmSingleFileBundle>true</WasmSingleFileBundle>
    <EventSourceSupport>false</EventSourceSupport>
    <UseSystemResourceKeys>true</UseSystemResourceKeys>
    <EnableUnsafeUTF7Encoding>false</EnableUnsafeUTF7Encoding>
    <HttpActivityPropagationSupport>false</HttpActivityPropagationSupport>
    <DebuggerSupport>false</DebuggerSupport>
  </PropertyGroup>
</Project>
```


WASM Results

Service	mean response time	90 percentile	99 percentile	max QPS	max Memory	Max CPU
c++=> WASM	9ms	12ms	16ms	1.2k	1 GB	15.4
c# => WASM	3 sec	3 sec	3 sec	1 (without the K)	1 GB	

Conclusion WASM

- c++ WASM can handle 1k QPS (file size 100 KB), c# WASM can handle 1 QPS (file size 30 MB)
- WASM doesn't seem like a real alternative, it currently needs to compile the file at each request

Efficient In memory caching

- Currently all data lookups from bidding script to KV state need to be serialized/de-serialized
- For real time bidding server must reply within 50ms
- **Optimizations here actually matter** (where we might not care much elsewhere)
- Network calls and IPC should be reduced or batched
- Over 30 caches in our Prod systems on PA POC


Efficient In memory caching, examples

- Efficient filtering, sending back all campaigns by country & all campaigns by domain and doing the intersection in the bidding layer
- ML Inference Sidecar, Inference side will require IPC, overhead will be significative, times 5 with internal ONX benchmarks

One option – inline datastructures

Inline large data structures for caching #62

 Open

 **Open** fhoering opened this issue on Jun 26 · 5 comments



fhoering commented on Jun 26 • edited ▾

Caching is an important mechanism for improving performance. I tried to inline a large datastructure to be able to reuse it in a readonly mode without querying the KV storage again. But using this the performance is getting very bad.

Can you have a look where the downlift is coming from ? Can it come from the large file sizes or from reinitializing the array all the time ?

```
const space = 10000000
const w = [0.004996137771039466,0.7571896910549433,0.7937806398622418,0.9647932057841067,0.00228668022937

function HandleRequest(executionMetadata, metadataKeys, signal) {
  let result = 0;
  for (let i = 0; i < 100; i++) {
    const index = Math.floor(Math.random() * space);
    result += w[index];
  }
  return result
}
};
```

Inlined datastructures

- JS engine doesn't handle inlined floats & structures in an efficient way
- Tested proposed option inlining ArrayBuffer with 1 mio elements

```
const buffer = new Uint8Array([1, 2, 3, 4]);
const floatArray = new Float32Array(buffer.buffer);

function HandleRequest(executionMetadata, ...udf_arguments) {
  for (let argument of udf_arguments) {
    let result = {};
    let key = argument.data;
    result[key] = floatArray[0];
    return result;
  }
};
```

Pure JS baseline

- getValues hook to access KV state
- Tested JS performance without getting state

```
function HandleRequest(executionMetadata, ...udf_arguments) {  
  for (let argument of udf_arguments) {  
    let key = argument.data;  
    const getValuesResult = JSON.parse(getValues(key));  
    return getValuesResult;  
  }  
};
```

In memory caching

Service	mean response time	90 percentile	99 percentile	max QPS	max Memory	Max CPU
key value service	6ms	8ms	9ms	5k	1 GB	15
kv service, inlined array	7ms	10ms	12ms	4.2k QPS	1 GB	15
kv service, no getvalues hook, (pure JS baseline)	7ms	10ms	12ms	6k QPS	1 GB	15

Conclusion in memory caching

- Didn't even try to deserialize a large object yet from `getValues`
- Pure JS baseline vs KV service => **+20%**
- We have 30 caches in production
- Explore Shared memory or Long running processes

Server Infra cost might quickly get out of control

JS ROMA vs Native c# => **x5**

No shared memory => **+50%** (30 caches in prod system)

ML inference side car => **x5** on ML inference (internal ONNX benchmark)

TEE and encrypted networking => **+20%** (usual symmetric encryption overhead)

Future work

More work is needed to mitigate potential infra cost increase

Shared memory or Inlining data structures

New ROMA BringYourOwn binary execution engine to allow custom languages beyond JS & WASM

Common method of benchmarking, improvements should be trackable easily, web load tests