

# Bidding and Auction Inference

# Agenda

- Background
- Overall flow
- Technical details
- Q&A

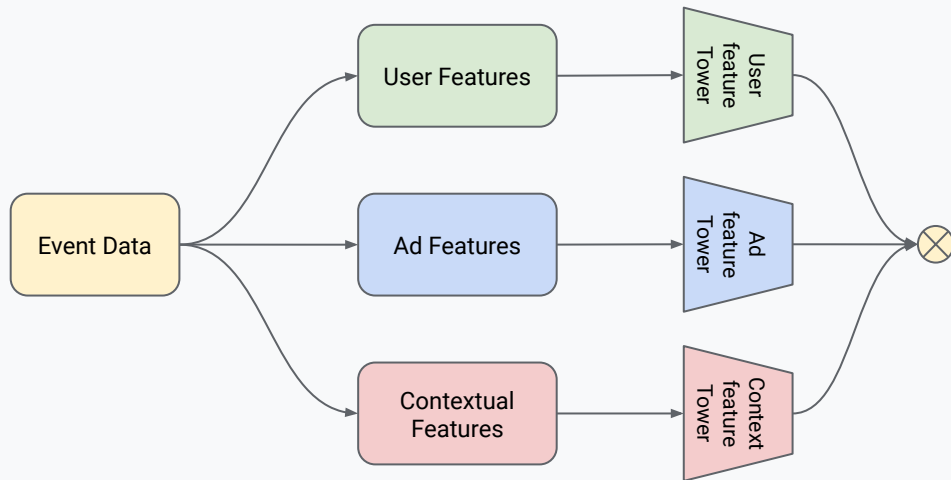
# Background

# Background

- Problem
  - Currently only models embedded inside JS / WASM are supported
  - Support first class capability to serve models inside TEE
- Adding Inference capability to Bidding and Auction services inside TEE
  - Initial support is for [Protected App Signals](#) on Android (Beta - EOQ2 2024)
  - Support for [Protected Audience](#) (Android, Chrome) to be added as a follow up (ETA TBD)
  - Initially supported on Bidding servers, can expand to others TEE servers in the future
- Scope for this presentation
  - Limited to serving models which are already trained
  - Does not cover model training
- Status
  - This is work in progress, seeking feedback from the larger community
  - Core pieces of implementation are available, but are not ready for production traffic

# Background

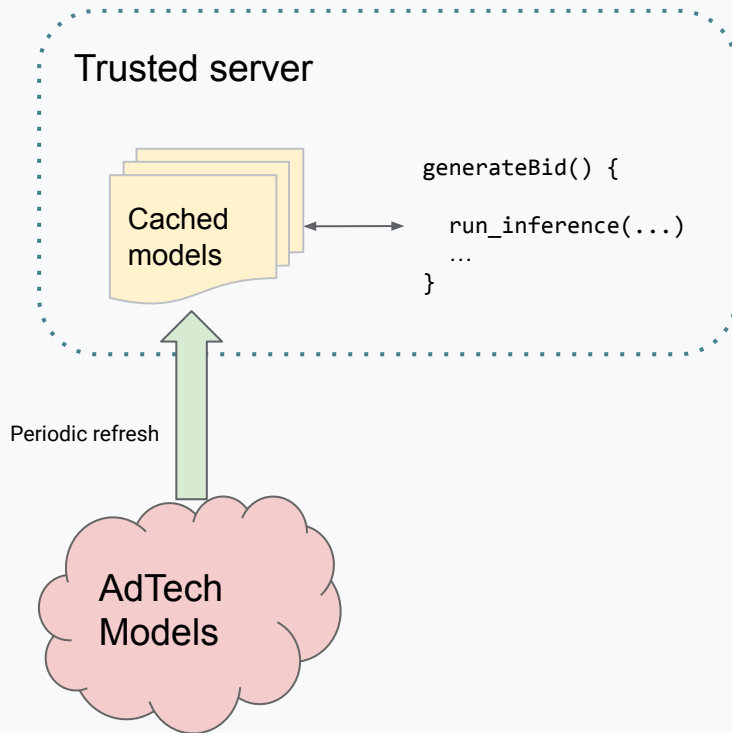
- Inference capability includes
  - Load models from cloud bucket
  - Call inference from JS worklets (e.g. `generateBid`, `prepareDataForAdRetrieval`)
  - Execute inference / predictions with standard backends (Tensorflow, PyTorch) inside TEE
- Support for factorized models
- For details, see [Inference explainer](#)



# Overall flow

# Executing models in B&A

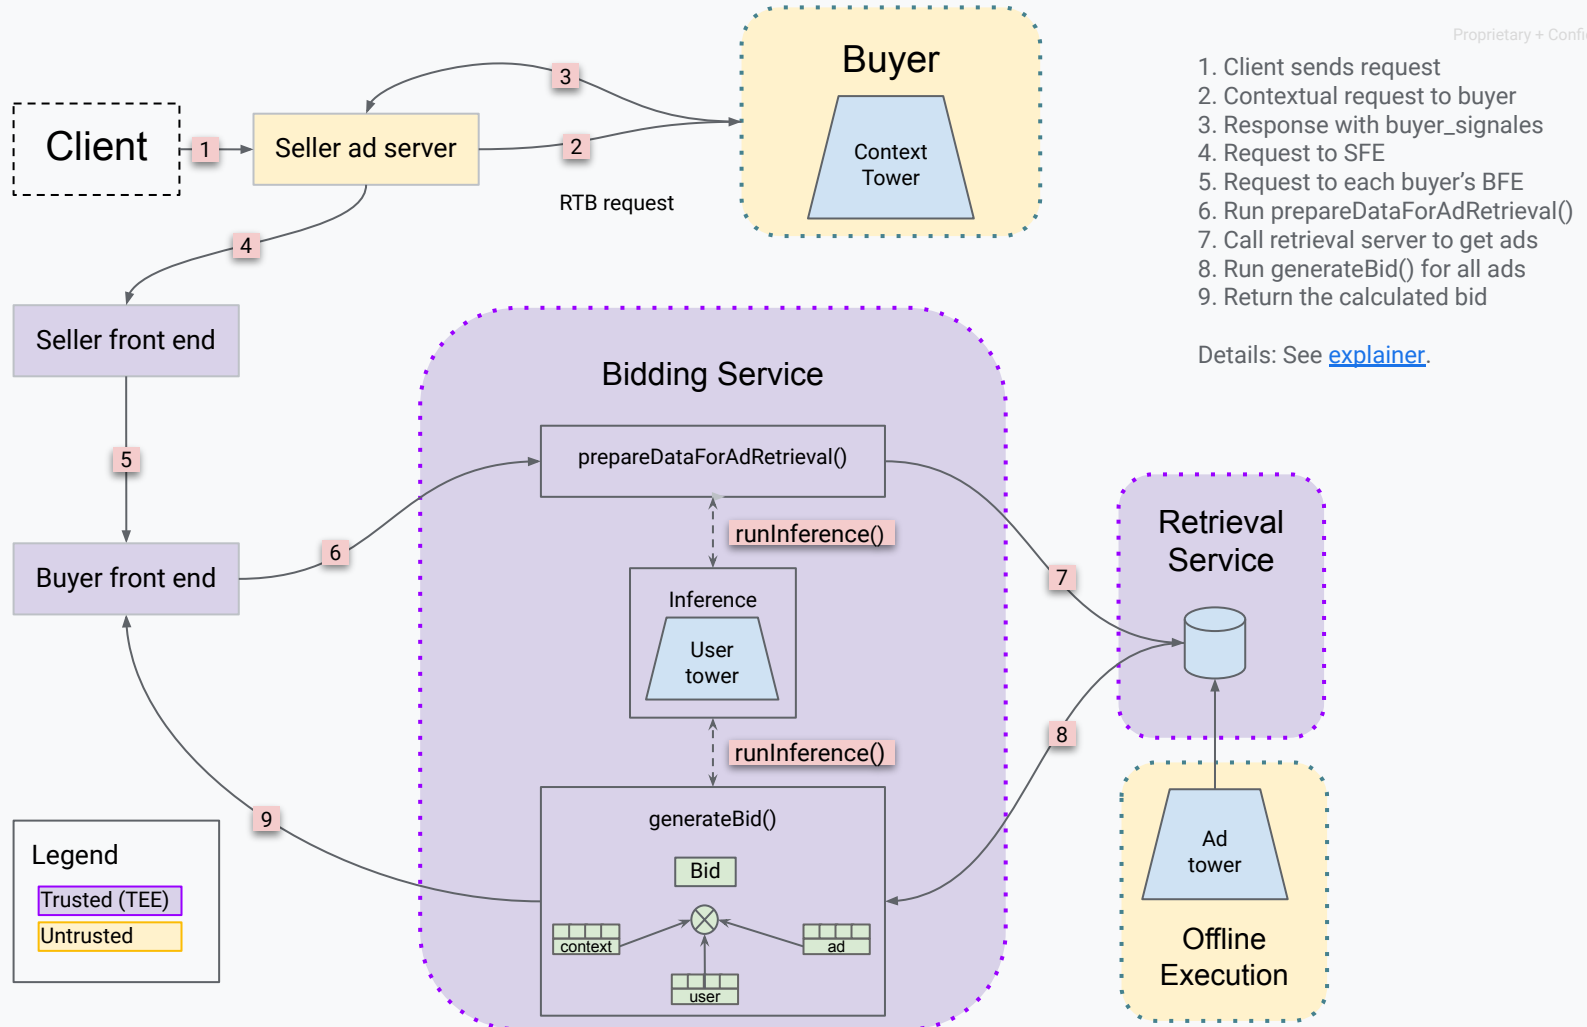
- Core functionality is to load and execute models
- This will support factorized and non-factorized models
- Models will be callable from JS
- Models will be periodically refreshed



# Factorized model flow

- B&A will support different plugin points for factorized models
  - B&A servers can do predictions with ML models inside TEE
  - RTB requests as a part of Contextual request (outside TEE) can generate contextual embeddings and pass to the `generateBid` function.
  - Pre-generated embeddings can be fetched from the retrieval server and be passed to `generateBid`. In this case, the model runs offline outside the TEE. The generated embeddings are uploaded to the retrieval server.
- Data flows for transferring embeddings to `generateBid` and consistent versioning of models and embeddings will be supported
- Ad-tech can use this support as building blocks to create flows for their inference





# Versioning

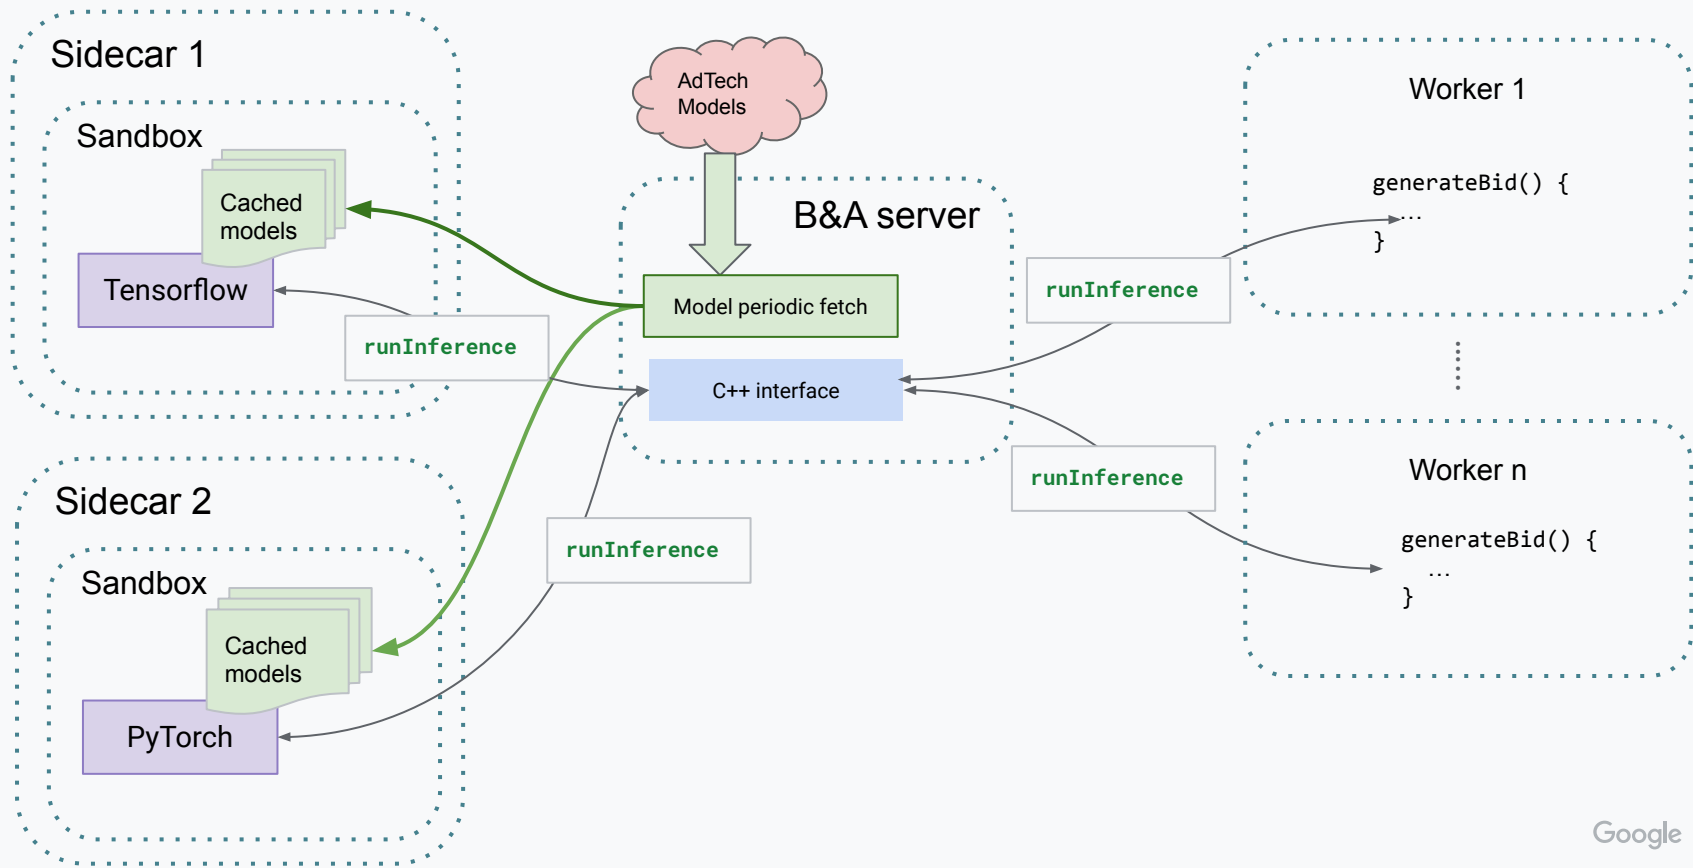
- Important to keep the versions of embedding and models in sync
- Passing version information supported between contextual path, retrieval server and the UDFs (e.g. `generateBid` and `prepareDataForAdRetrieval`)
- Specific versions of embeddings can be coordinated through this version information.
- Model versions can be encoded in the path.
  - Path is the main primitive for identifying a model identifier
  - The path can encode the name and version information together
  - Paths of all loaded models are made available to the UDFs, so the right version of the model can be selected
- Details: [Versioning section](#) in the Inference explainer

# Technical details

# Model execution

- Two main ways of executing models inside TEE
  - Embedded models
  - Externally provided Models for supported backends (initially Tensorflow and PyTorch)
- Embedded models are fully contained inside the ad tech JS / WASM
- Non-embedded models for supported platforms will be executed by the B&A code by linking C++ libraries for each model type
- Initially Tensorflow and PyTorch (Libtorch) are supported. Support for other systems (e.g. ONNX) can be added later based on ad tech demand.
- Initially models will run on CPU. As accelerators become available we will explore integration - timelines are TBD

# Sidecar model execution



# API

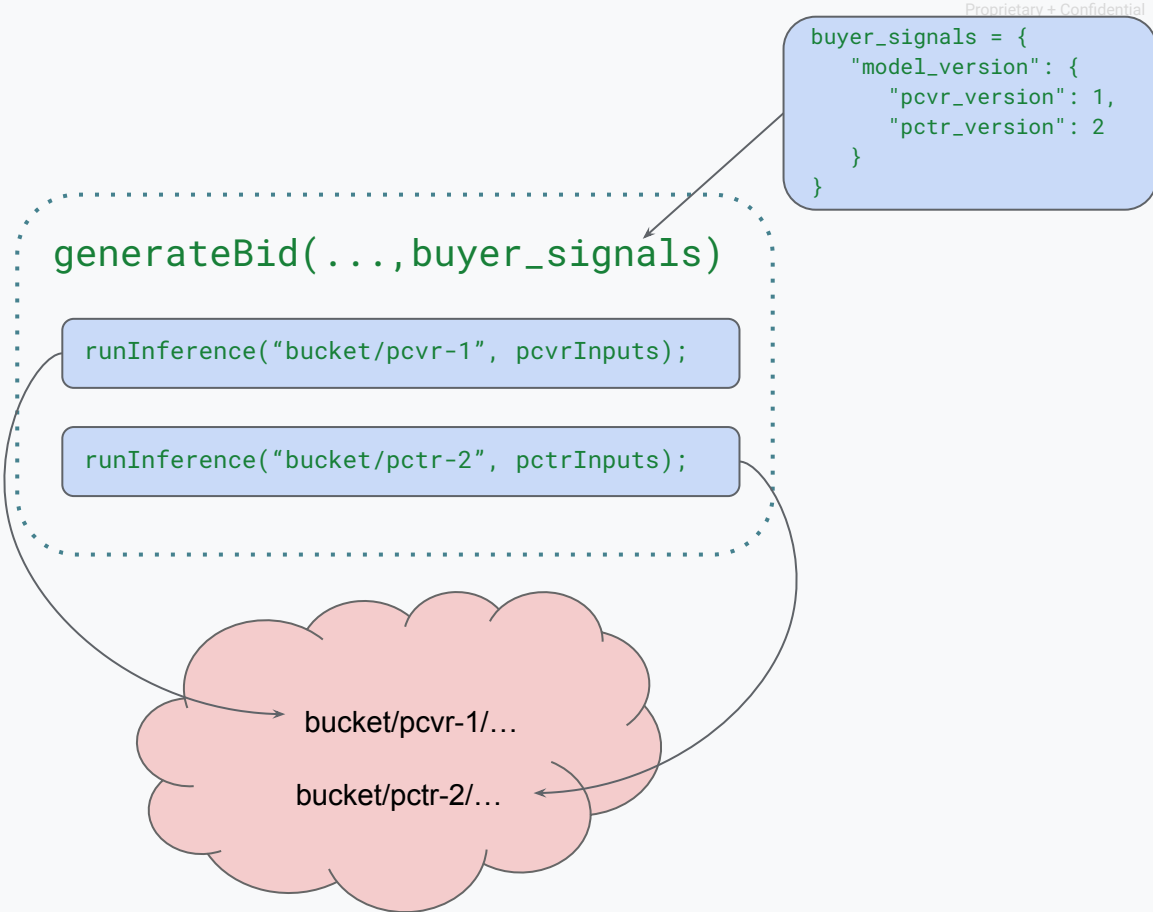
- `runInference(request)` is the main way of calling inference.
- Batch API supports calling multiple models in the same request
- For each model, inputs and outputs can be arbitrary tensors.

## Request for one model:

```
"model_path": "my_bucket/models/pcvr/1/",  
"tensors": [{  
  "tensor_name": "user_signal_1",  
  "data_type": "DOUBLE",  
  "tensor_shape": [2, 1],  
  "tensor_content": ["0.454920", "-0.25752"]  
}]
```

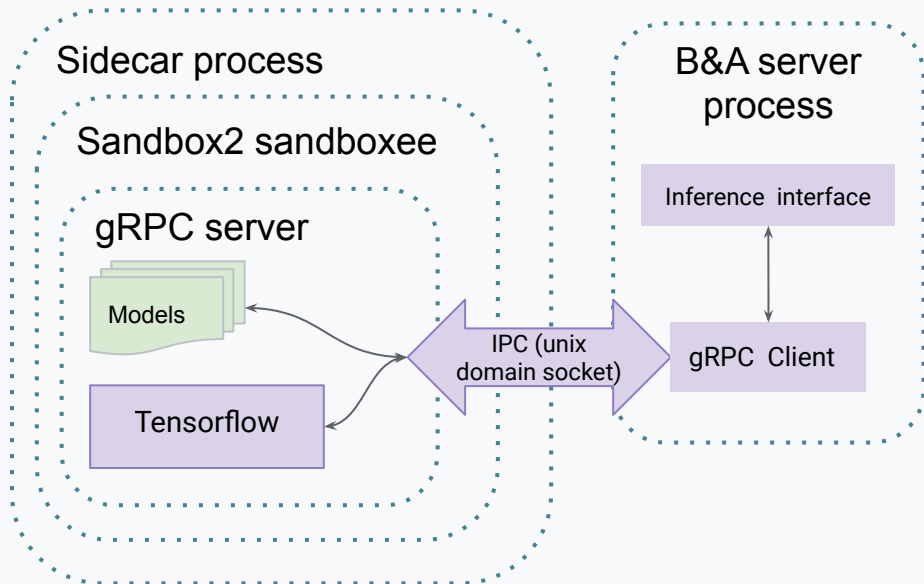
# Model loading

- Models will be loaded from cloud buckets
- Once models are loaded, they can be used for inference
- The main way to address a model is by its path
- Versioning info can be used inside `generateBid` (and other UDFs) to create the correct model path.
- There is an API that can be used from `generateBid` to get a list of loaded model paths.



# Sandboxing

- Inference runs in a separate process called the inference sidecar
- The sidecar is a separate binary compiled with one backend execution system (e.g. a specific version of Tensorflow).
- The ML backend is sandboxed using [Sandbox2](#).
- Model loading and inference happens only inside the sandbox
- Inside the sandbox, we run a gRPC server that receives requests and responds to the host binary (e.g. Bidding server).
- gRPC Communication between the host and the sidecar happen through a unix domain socket managed by Sandbox2.





# Privacy mitigations

- Privacy mitigation are still in progress and will be put in place before MVP
- Models and ML backends are completely sandboxed
- Models will periodically be reset to their original state to make sure they are not stateful
- More details will be added to the explainer in the future

# Q&A

# Feedback themes

- Preferred ML backends
- Model size, scale, performance
- Model versions and update frequency
- Operational and quality metrics
- Cost and Utility
- Detailed questionnaire: [this doc](#)
- Feedback: <https://github.com/WICG/protected-auction-services-discussion/issues/59>

05

# Appendix

# References

- [Inference explainer](#)
- [Feedback questionnaire](#)
- [Bidding and Auction services explainer](#)
- [Protected app signals explainer](#)
- [Protected Auctions explainer](#)